

▼ Classification Using Tensorflow

Richard Kuo, 201807016 ver. 0.1.1 AdamOptimizer

Code borrowed from:

- [Cifar-10 Classification using Keras Tutorial](#)
- [Object Recognition with Convolutional Neural Networks in the Keras Deep Learning Library](#)
- [Convolutional Neural Networks \(CNN\) for CIFAR-10 Dataset](#)
- [Deep-math-machine-learning.ai](#)
- [Keras code example](#)

▼ Import libraries

```
import time
import sys
import os

from __future__ import print_function

from keras.utils import print_summary, to_categorical, np_utils

# Import Tensorflow with multiprocessing
import tensorflow as tf
import multiprocessing as mp

# Loading the CIFAR-10 datasets
from keras.datasets import cifar10

# Plot
import matplotlib.pyplot as plt
% matplotlib inline
from scipy.misc import toimage

import numpy as np
np.random.seed(2018)
```

☞ Using TensorFlow backend.

▼ Define constants

This will make code more flexible.

```
# Declare variables
```

▼ Load and display dataset

After data loading, to verify and better understand the dataset; sample some them. For more complicate dataset, plot, explore the contents.

- shapes
- sizes
- sample values

```
# load data, instead of using the built-in function, this can be done with pyth
(X_train, Y_train), (X_test, Y_test) = cifar10.load_data()

X_train # tensor type
Y_train
print('X_train shape:', X_train.shape)
print('Y_train shape:', Y_train.shape)
print(X_train.shape[0], 'train samples')
print(X_test.shape[0], 'test samples')

print("Value of the first element of X_train:")
print(X_train[0])
print("Value of the first element of Y_train:")
print(Y_train[0])

# create a grid of 3x3 images
print("X can be converted back to original images via utility function:")
for i in range(0, 9):
    plt.subplot(330 + 1 + i)
    plt.imshow(toimage(X_train[i]))
# show the plot
plt.show()
```



```

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.ta
170500096/170498071 [=====] - 21s 0us/step
X_train shape: (50000, 32, 32, 3)
Y_train shape: (50000, 1)
50000 train samples
10000 test samples
Value of the first element of X_train:
[[ 59  62  63]
 [ 43  46  45]
 [ 50  48  43]
 ...
 [158 132 108]
 [152 125 102]
 [148 124 103]]

[[ 16  20  20]
 [  0   0   0]
 [ 18   8   0]
 ...
 [123  88  55]
 [119  83  50]
 [122  87  57]]

[[ 25  24  21]
 [ 16   7   0]
 [ 49  27   8]
 ...
 [118  84  50]
 [120  84  50]
 [109  73  42]]

...

[[208 170  96]
 [201 153  34]
 [198 161  26]
 ...
 [160 133  70]
 [ 56  31   7]
 [ 53  34  20]]

[[180 139  96]

```

There are three different sets/formats: python, Matlab and binary. Binary format has broke the dataset into smaller size so we can do batch process.

For python version:

data -- a 10000x3072 numpy array of uint8s. Each row of the array stores a 32x32 (=1024) colour image. The first 1024 entries contain the red channel values, the next 1024 the green, and the final 1024 the blue. The image is stored in row-major order, so that the first 32 entries of the array are the red channel values of the first row of the image, [59 43 50 ... 158 152 148].

The color of the first pixel of the image is R = 59, G= 62, B = 63.

Each image is represented by 3 X 32 X 32 (colors, width and height) tensor. The first diamension is for Red, the next sheet is for Blue, then the next sheet is for Green; like 3 sheets of paper, which are

size 32 by 32 grids. Each grid records the values of the color, RGB, of image values (0–255). There are 50000 images for X_train.

Tensor is data structure, for example,

```
1D  [1,2,3,4]          # one row
2D  [                  # one sheet with two rows
      [1,2,3,4],
      [5,6,7,8]]
3D  [                  # three sheets, each sheet represents some
type of measurement
      [                  # for example, people's height, or
intensity of color red
      [1,2,3,4],
      [5,6,7,8]],
      [                  # for example, people's weight, or
intensity of color blue, we only have two colors.
      [7,2,3,4],
      [6,6,7,8]]]
```

labels -- a list of 10000 numbers in the range 0-9. The number at index i indicates the label of the ith image in the array data. For example, label_names[0] == "airplane", label_names[1] == "automobile", etc.

label_names = ['airplane','automobile','bird','cat','deer', 'dog','frog','horse','ship','truck']

Y_train and Y_test are a vector (1D tensor). Y_train[0] = [6] means the label of the first element is a frog.

X_train and X_test are 3D tensor, which can be converted back to image.



See [toronto data repo](#) for more details.

▼ Pre-process Data

The data may need to be pre-processed to fit ML library we want to use.

Normalize

RGB value is range from 0-255. It would be easier to work from 0-1.

```
# print(X_train[0])
X_train = X_train.astype('float32')
X_test  = X_test.astype('float32')
X_train /= 255.0
X_test  /= 255.0

X_train[0]
```



```

array([[0.23137255, 0.24313726, 0.24705882],
       [0.16862746, 0.18039216, 0.1764706 ],
       [0.19607843, 0.1882353 , 0.16862746],
       ...,
       [0.61960787, 0.5176471 , 0.42352942],
       [0.59607846, 0.49019608, 0.4          ],
       [0.5803922 , 0.4862745 , 0.40392157]],

       [[0.0627451 , 0.07843138, 0.07843138],
       [0.          , 0.          , 0.          ],
       [0.07058824, 0.03137255, 0.          ],
       ...,
       [0.48235294, 0.34509805, 0.21568628],
       [0.46666667, 0.3254902 , 0.19607843],
       [0.47843137, 0.34117648, 0.22352941]],

       [[0.09803922, 0.09411765, 0.08235294],
       [0.0627451 , 0.02745098, 0.          ],
       [0.19215687, 0.10588235, 0.03137255],
       ...,
       [0.4627451 , 0.32941177, 0.19607843],
       [0.47058824, 0.32941177, 0.19607843],
       [0.42745098, 0.28627452, 0.16470589]],

       ...,

       [[0.8156863 , 0.6666667 , 0.3764706 ],
       [0.7882353 , 0.6          , 0.13333334],
       [0.7764706 , 0.6313726 , 0.10196079],
       ...,
       [0.627451 , 0.52156866, 0.27450982],
       [0.21960784, 0.12156863, 0.02745098],
       [0.20784314, 0.13333334, 0.07843138]],

       [[0.7058824 , 0.54509807, 0.3764706 ],
       [0.6784314 , 0.48235294, 0.16470589],
       [0.7294118 , 0.5647059 , 0.11764706],
       ...,
       [0.72156864, 0.5803922 , 0.36862746],
       [0.38039216, 0.24313726, 0.13333334],
       [0.3254902 , 0.20784314, 0.13333334]],

       [[0.69411767, 0.5647059 , 0.45490196],
       [0.65882355, 0.5058824 , 0.36862746],
       [0.7019608 , 0.5568628 , 0.34117648],
       ...,

```

The value of the first pixel R-channel is $59/255 = 0.2313$, which is the current X_{train} value divided by 255. If we run this cell again, the value will be 0.00090734, which is $= 0.2313/255$.

The, we need to see what is train label Y_{train} looks like.

Y_{train}



```
array([[6],
       [9],
       [9],
       ...,

```

It is a 50000 integer array.

```
print(Y_train[0], Y_train[49999])
```

```
[6] [1]
```

We need to convert Y_train to one_hot array, from 50000 x 1 to 50000 x 10, for example, the first row from [6] to [0,0,0,0,0,1,0,0,0,0]. We can use a keras' utility `keras.utils.np_utils.to_categorical`, see discussion about `to_categorical`.

```
# create one hot vector
```

```
Y_train = np_utils.to_categorical(Y_train)
Y_test = np_utils.to_categorical(Y_test)
```

```
Y_train
```

```
[0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 1.],
 [0., 0., 0., ..., 0., 0., 1.],
 ...,
 [0., 0., 0., ..., 0., 0., 1.],
 [0., 1., 0., ..., 0., 0., 0.],
 [0., 1., 0., ..., 0., 0., 0.]], dtype=float32)
```

```
Y_train[0]
```

```
[0., 0., 0., 0., 0., 0., 1., 0., 0., 0.], dtype=float32)
```

▼ Tensorflow

We will try to develop the solution with two different frameworks: Tensorflow and Keras, do Tensorflow first. The architecture looks like

[input (X)] -> [convolution (C1)] -> [pooling (P1)] -> [convolution (C2)] -> [pooling (P2)] -> [convolution (C3)] -> [pooling (P3)] -> [FC] -> [softmax] -> [output (Y predict)]

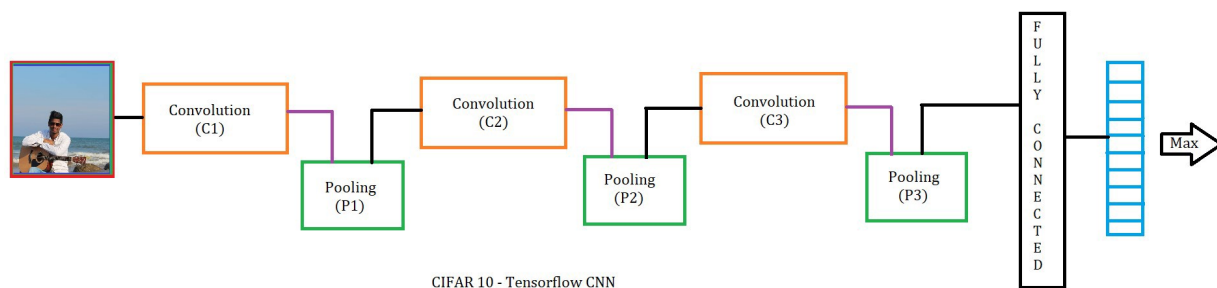


Figure from <https://medium.com/deep-math-machine-learning-ai/chapter-8-1-code-for-convolutional-neural-networks-tensorflow-and-keras-theano-33bef285dd93>

▼ Model

▼ placeholder

Inserts a placeholder for a tensor that will be always fed, see [doc].

(https://www.tensorflow.org/api_docs/python/tf/placeholder)

```
tf.placeholder(  
    dtype,  
    shape=None,  
    name=None  
)
```

Args:

- dtype: The type of elements in the tensor to be fed.
- shape: The shape of the tensor to be fed (optional). If the shape is not specified, you can feed a tensor of any shape.
- name: A name for the operation (optional).

The difference is that with `tf.Variable` you have to provide an initial value when you declare it. With `tf.placeholder` you don't have to provide an initial value and you can specify it at run time with the `feed_dict` argument inside `Session.run`

```
import tensorflow as tf  
  
X = tf.placeholder("float", [None, 32, 32, 3])  
Y = tf.placeholder("float", [None, 10])
```

Weight

define a function to populate specific shape of tensor initially with random numbers.

`tf.random_normal` returns a tensor of specific shape with random number, see [doc](#)

```
tf.random_normal(  
    shape,  
    mean=0.0,  
    stddev=1.0,  
    dtype=tf.float32,  
    seed=None,  
    name=None  
)
```

Args:

- shape: A 1-D integer Tensor or Python array. The shape of the output tensor.
- mean: A 0-D Tensor or Python value of type dtype. The mean of the normal distribution.
- stddev: A 0-D Tensor or Python value of type dtype. The standard deviation of the normal distribution.
- dtype: The type of the output.
- seed: A Python integer. Used to create a random seed for the distribution. See `tf.set_random_seed` for behavior.
- name: A name for the operation (optional).

▼ Layer

tf's Layers (contrib) package

- provide ops that take care of creating variables that are used internally in a consistent way,
- provide the building blocks for many common machine learning algorithms.

```
def init_weights(shape):
    return tf.Variable(tf.random_normal(shape, stddev=0.01))

W_C1 = init_weights([3,3,3,32])      # 3x3x3 conv, 32 outputs
W_C2 = init_weights([3,3,32,64])     # 3x3x32 conv, 64 outputs
W_C3 = init_weights([3, 3, 64, 128]) # 3x3x64 conv, 128 outputs

W_FC = init_weights([128 * 4 * 4, 625]) # FC 128 * 4 * 4 inputs, 625 outputs
W_O  = init_weights([625, 10])        # FC 625 inputs, 10 outputs (labels)

p_keep_conv = tf.placeholder("float") # for dropouts as percentage
p_keep_hidden = tf.placeholder("float")
```

Convolution

tf's Neural Network (nn) API supports (partially excerpt here):

- Convolution ops sweep a 2-D filter over a batch of images, applying the filter to each window of each image of the appropriate size.
- Activation Function ops provide various types of nonlinearities for use in neural networks.
- Pooling ops sweep a rectangular window over the input tensor, computing a reduction operation for each window (average, max, or max with argmax). ... for details see [nn](#)

`tf.nn.conv2d` computes a 2-D convolution given 4-D input and filter tensors, see [tf.conv2d](#) for details.

```
tf.nn.conv2d(
    input,
    filter,
    strides,
    padding,
    use_cudnn_on_gpu=True,
    data_format='NHWC',
    dilations=[1, 1, 1, 1],
    name=None
)
```


Activation

ReLU is one of activation function supported in tf. `tf.nn.relu` takes a tensor (type options: float32, float64, int32, uint8, int16, int8, int64, bfloat16, uint16, half, uint32, uint64.) as input, rectifies it, and returns a tensor with the same type as input.

```
tf.nn.relu(  
    features,  
    name=None  
)
```

Pooling

tf supports several different pooling ops:

- `tf.nn.avg_pool`
- `tf.nn.max_pool`
- `tf.nn.max_pool_with_argmax`
- `tf.nn.avg_pool3d`
- `tf.nn.max_pool3d`
- `tf.nn.fractional_avg_pool`
- `tf.nn.fractional_max_pool`
- `tf.nn.pool`

Max_pooling is a very popular pooling.

```
tf.nn.max_pool(  
    value,  
    ksize,  
    strides,  
    padding,  
    data_format='NHWC',  
    name=None  
)
```

Args:

- `value`: A 4-D Tensor of the format specified by `data_format`.
- `ksize`: A 1-D int Tensor of 4 elements. The size of the window for each dimension of the input tensor.
- `strides`: A 1-D int Tensor of 4 elements. The stride of the sliding window for each dimension of the input tensor.
- `padding`: A string, either 'VALID' or 'SAME'. The padding algorithm. See the comment here
- `data_format`: A string. 'NHWC', 'NCHW' and 'NCHW_VECT_C' are supported.
- `name`: Optional name for the operation.

Returns: A Tensor of format specified by `data_format`. The max pooled output tensor.

▼ Dropout

`tf.nn.dropout` computes dropout. The dropout is used to prevent overfitting. Some % of neurons will be randomly dropped out (don't get computed).

With probability `keep_prob`, outputs the input element scaled up by $1 / \text{keep_prob}$, otherwise outputs 0. The scaling is so that the expected sum is unchanged, for details see [dropout](#).

```
tf.nn.dropout(
    x,
    keep_prob,
    noise_shape=None,
    seed=None,
    name=None
)

def model(X, W_C1, W_C2, W_C3, W_FC, W_O, p_keep_conv, p_keep_hidden):

    C1 = tf.nn.relu(tf.nn.conv2d(X, W_C1, strides=[1,1,1,1], padding = "SAME"))
    P1 = tf.nn.max_pool(C1, ksize=[1,2,2,1], strides=[1,2,2,1], padding = "SAME")
    D1 = tf.nn.dropout(P1, p_keep_conv) # 1st dropout at conv

    C2 = tf.nn.relu(tf.nn.conv2d(D1, W_C2, strides=[1,1,1,1], padding = "SAME"))
    P2 = tf.nn.max_pool(C2, ksize=[1,2,2,1], strides=[1,2,2,1], padding = "SAME")
    D2 = tf.nn.dropout(P2, p_keep_conv) # 2nd dropout at conv

    C3 = tf.nn.relu(tf.nn.conv2d(D2, W_C3, strides=[1,1,1,1], padding = "SAME"))
    P3 = tf.nn.max_pool(C3, ksize=[1,2,2,1], strides=[1,2,2,1], padding = "SAME")

    P3 = tf.reshape(P3, [-1, W_FC.get_shape().as_list()[0]]) # reshape to (?)
    D3 = tf.nn.dropout(P3, p_keep_conv) # 3rd dropout at conv

    FC = tf.nn.relu(tf.matmul(D3, W_FC))
    FC = tf.nn.dropout(FC, p_keep_hidden) # dropout at fc

    output = tf.matmul(FC, W_O)

    return output
```

▼ Cost function, Optimizer and Predicted function.

todo - read [Neural Network Optimization Algorithms](#)

```
Y_pred = model(X, W_C1, W_C2, W_C3, W_FC, W_O, p_keep_conv, p_keep_hidden) #
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits_v2(logits = Y_pre

optimizer = tf.train.AdamOptimizer(0.001, 0.9).minimize(cost)
# optimizer = tf.train.RMSPropOptimizer(0.001, 0.9).minimize(cost)

predict_op = tf.argmax(Y_pred, 1)
```

▼ Session

To run TensorFlow model, we need to create a session and feed the real data.

```

# reshape input data per tf
X_train = X_train.reshape(-1,32,32,3)
X_test = X_test.reshape(-1,32,32,3)

epochs = 50
import numpy as np
with tf.Session() as sess:
    # you need to initialize all variables
    sess.run(tf.global_variables_initializer())

    for epoch in range(epochs):

        for start, end in zip(range(0, len(X_train), 128), range(128, len(X_train), 128)):
            sess.run(optimizer, feed_dict={X : X_train[start:end] , Y : Y_train[start:end],
                                            p_keep_conv: 0.8, p_keep_hidden: 0.5})

        if epoch % 10 == 0:
            accuracy= np.mean(np.argmax(Y_test, axis=1) ==
                               sess.run(predict_op, feed_dict={X: X_test, p_keep_conv: 0.8, p_keep_hidden: 0.5}))

            print("epoch : {} and accuracy : {}".format(epoch, accuracy))

            print("testing labels for test data")
            print(sess.run(predict_op, feed_dict={X: X_test, p_keep_conv: 1.0, p_keep_hidden: 0.5}))

        print("Final accuracy : {}".format(np.mean(np.argmax(Y_test, axis=1) ==
                                                    sess.run(predict_op, feed_dict={X: X_test, p_keep_conv: 0.8, p_keep_hidden: 0.5}))))

```

```

☞ epoch : 0 and accuracy : 0.4395
testing labels for test data
[3 8 8 ... 5 6 7]
epoch : 10 and accuracy : 0.7247
testing labels for test data
[3 8 8 ... 5 1 7]
epoch : 20 and accuracy : 0.7516
testing labels for test data
[5 8 0 ... 5 1 7]
epoch : 30 and accuracy : 0.7592
testing labels for test data
[5 8 0 ... 5 1 7]
epoch : 40 and accuracy : 0.7673
testing labels for test data
[3 8 8 ... 5 1 7]
Final accuracy : 0.7587

```

